

# Data center emulation and traffic management via a Cloud SDN Controller

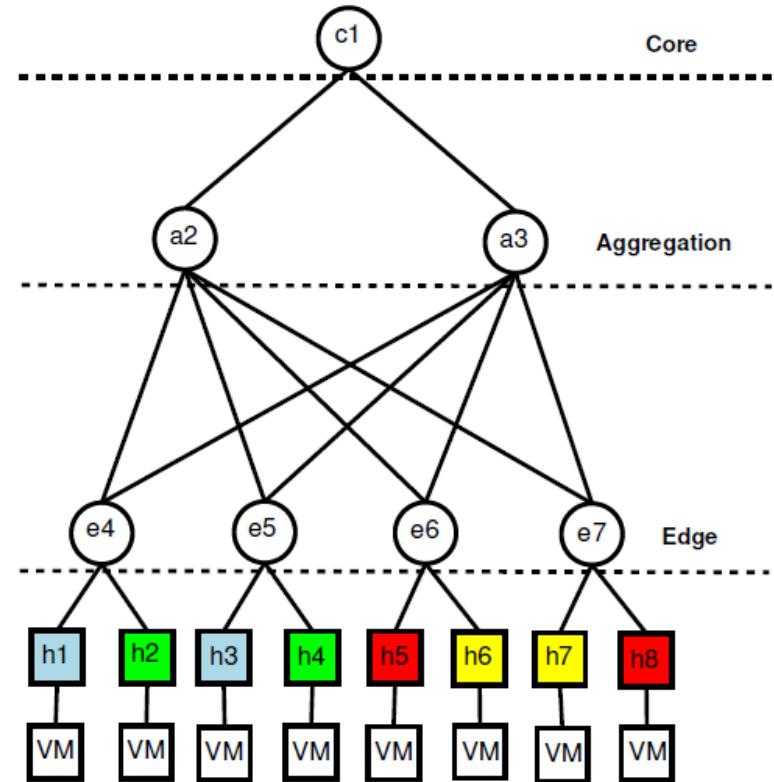
Emmanouil Lakiotakis

Course: HY436

Fall 2020

# Definition

- Data center: a facility used to host computer systems and associated components, such as telecommunications and storage systems.
- In practice, a number of Virtual Machines hosted by distinct end-host servers. Servers are grouped based on their services creating **tenant groups** (different colors).



# Data center requirements

- Efficient traffic routing across its expensive fabric
- Isolation
- Proper traffic engineering
- VM migration

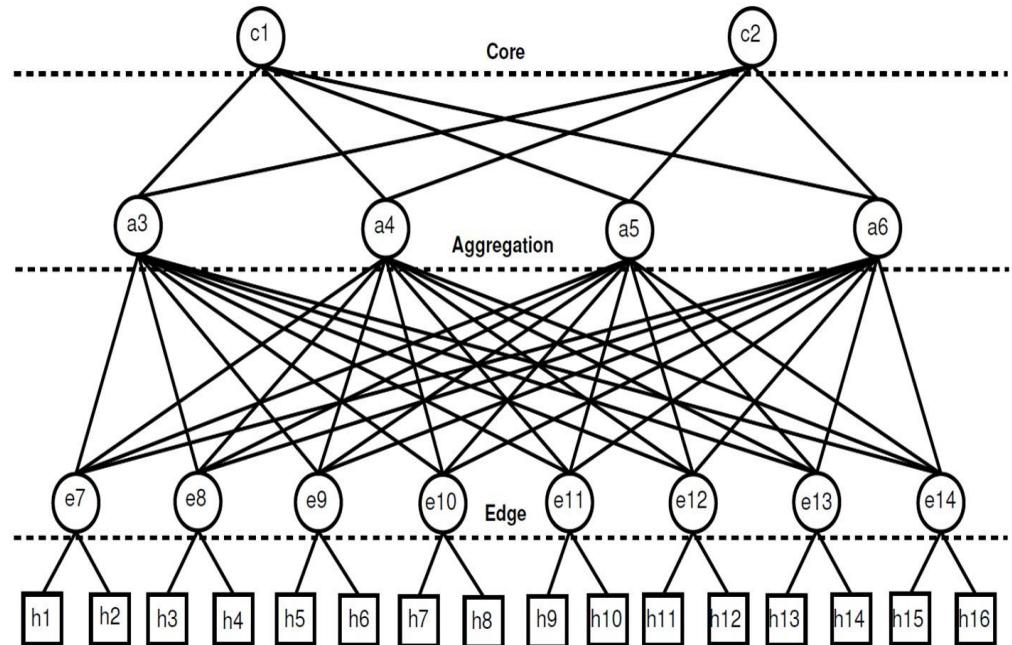
# Assignment objectives

- Emulate a ‘sample’ data center topology via **Mininet**
- Satisfy conventional data center networking requirements using basic SDN principles via **POX**

# Step 1: Network Topology

- Clos topology: tree-like with 3 layers
  - Core switches
  - Aggregation switches
  - Edge/access switches
- 2 parameters:
  - $c$ : # core switches
  - $f$ : (fanout) # children switches per parent
- **Follow depicted notation**

Example topology:  $c=2, f=2$

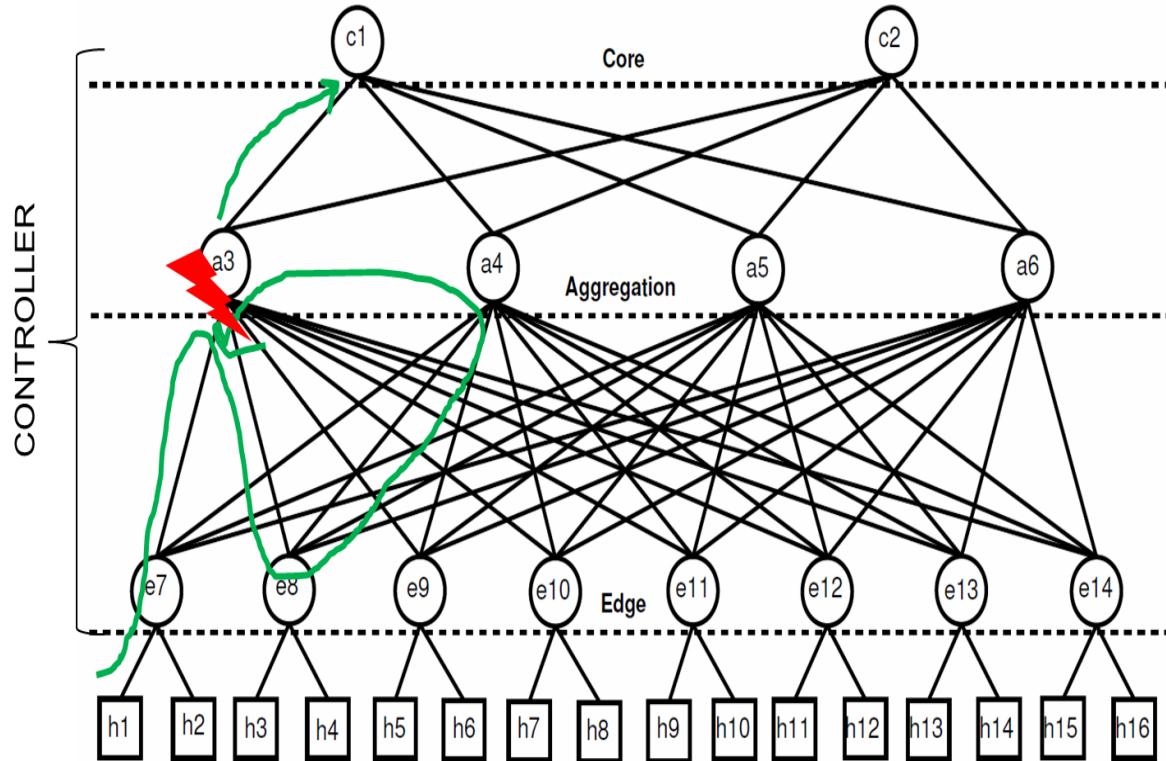


# Step 2: Traffic Routing

- Based on shortest path routing between hosts
- Traffic classification/differentiation based on different protocols (similar to ToS matching)
- **Real traffic engineering objectives used in conventional data centers**

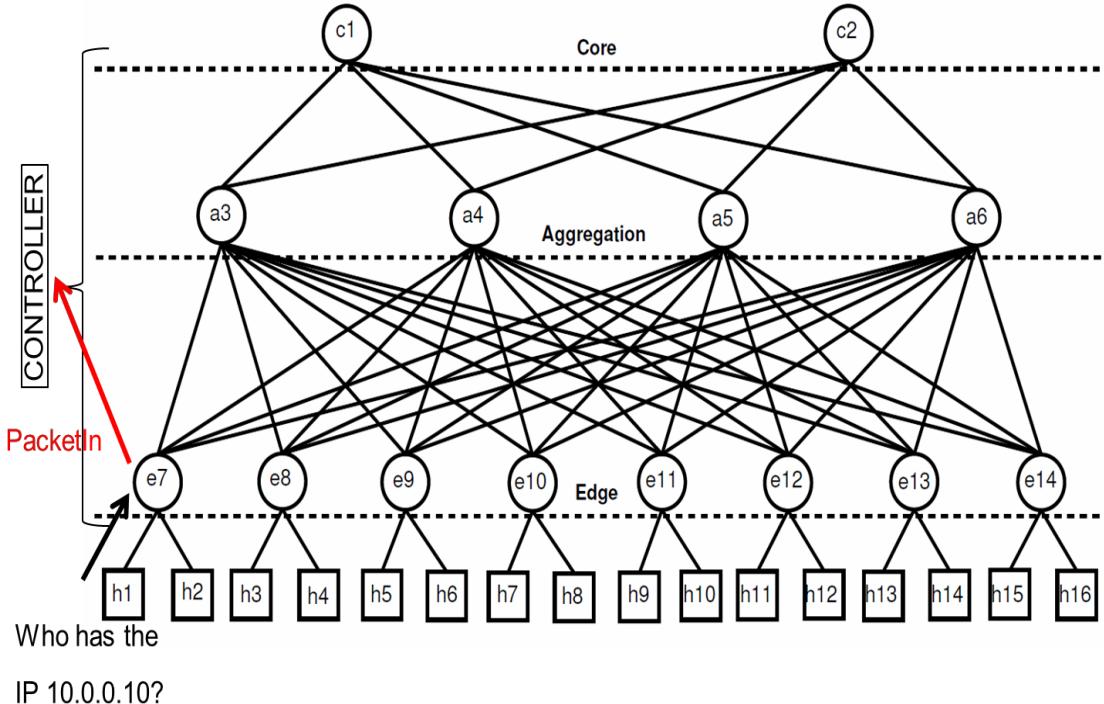
# ARP packet handling

- In large scale topologies, conventional flooding approach can lead to routing loops
- Avoid ARP forwarding to higher switch levels
- **Solution: ARP proxying**



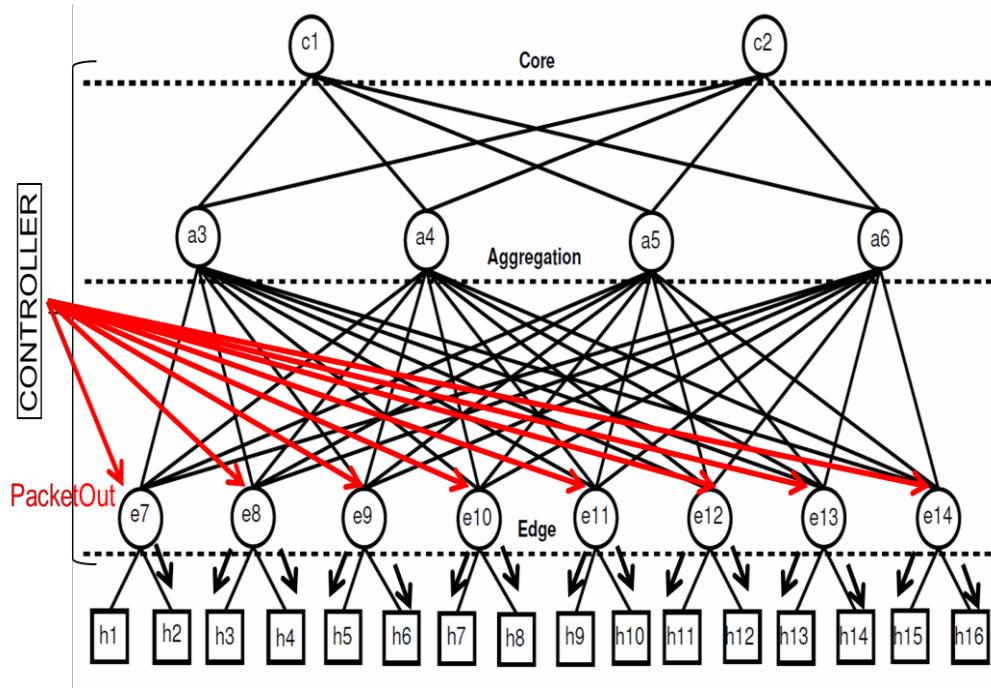
# ARP proxying: ARP requests

- ARP requests raise `PacketIn` events
- 2 possible cases:
  - The Controller **does not know** how to answer to the ARP request
  - The controller **knows** how to answer the ARP request



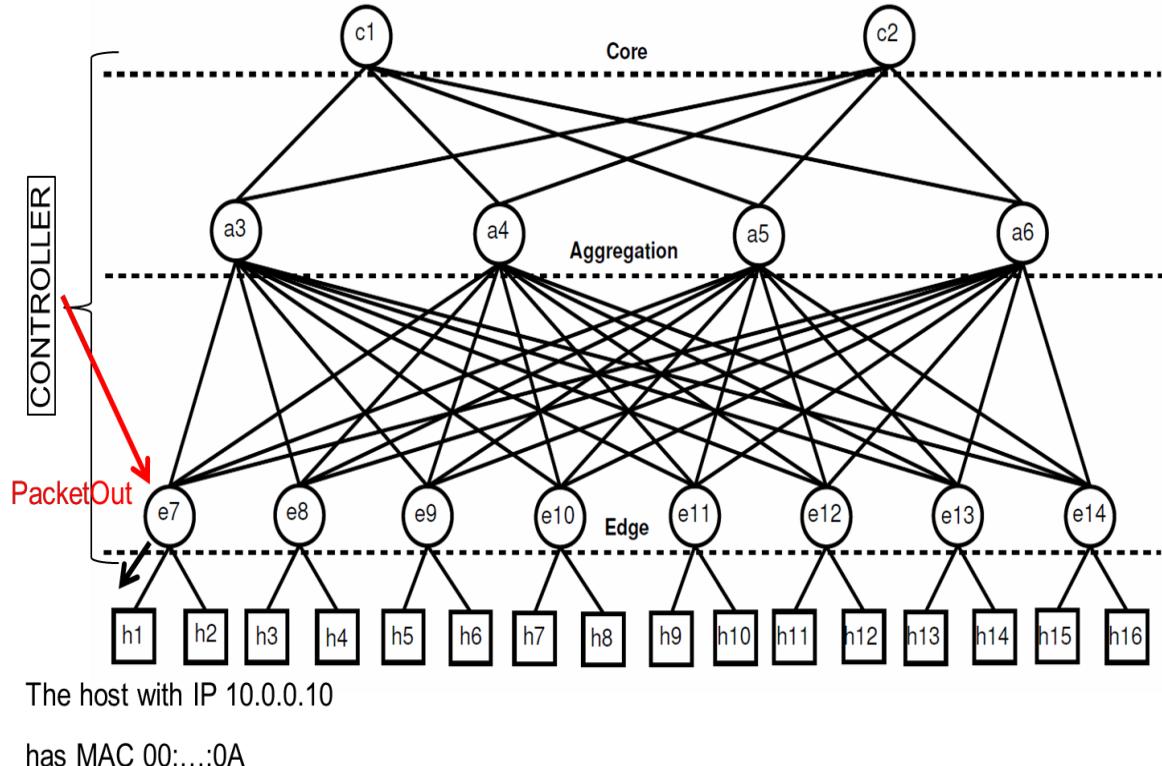
# ARP proxying: ARP requests (a)

- When the controller **does not** know how to answer to ARP request:
  - 'Sophisticated' flooding only to ports that lead to hosts (`flood_on_switch_edge()`)
  - Avoid routing loops
  - When host replies, forward the reply to the requesting host
- The controller knows which ports lead to switches and which ports lead to hosts



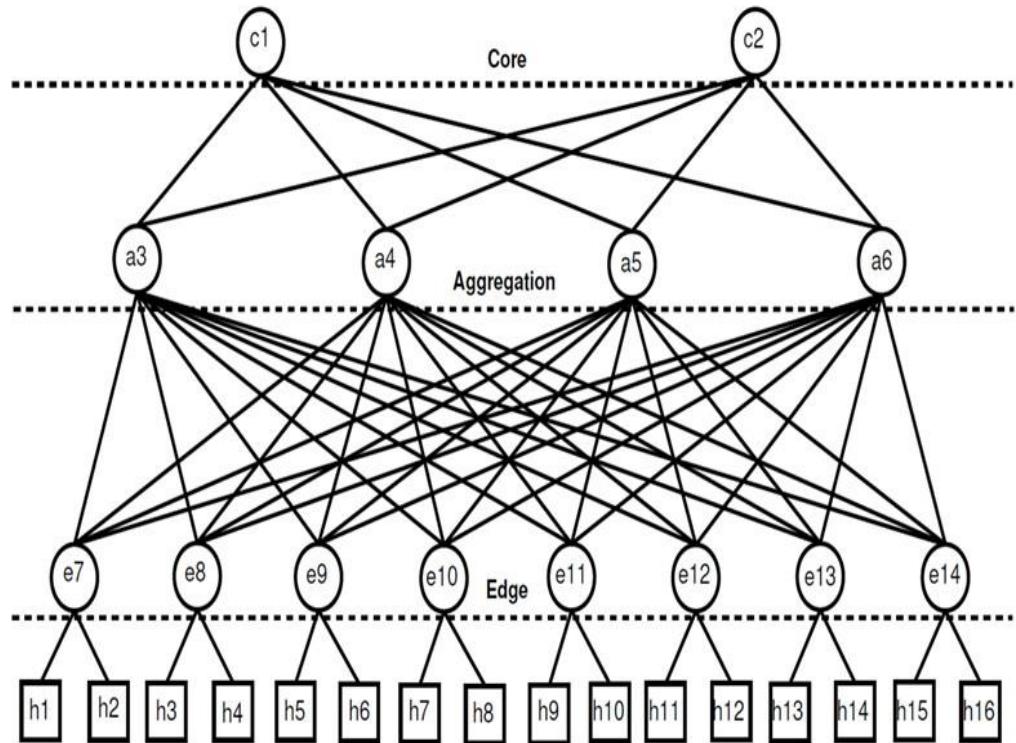
# ARP proxying: ARP requests (b)

- When the controller **knows** how to answer to ARP request
- Acts as proxy creating the appropriate ARP reply and forwarding to the requesting host



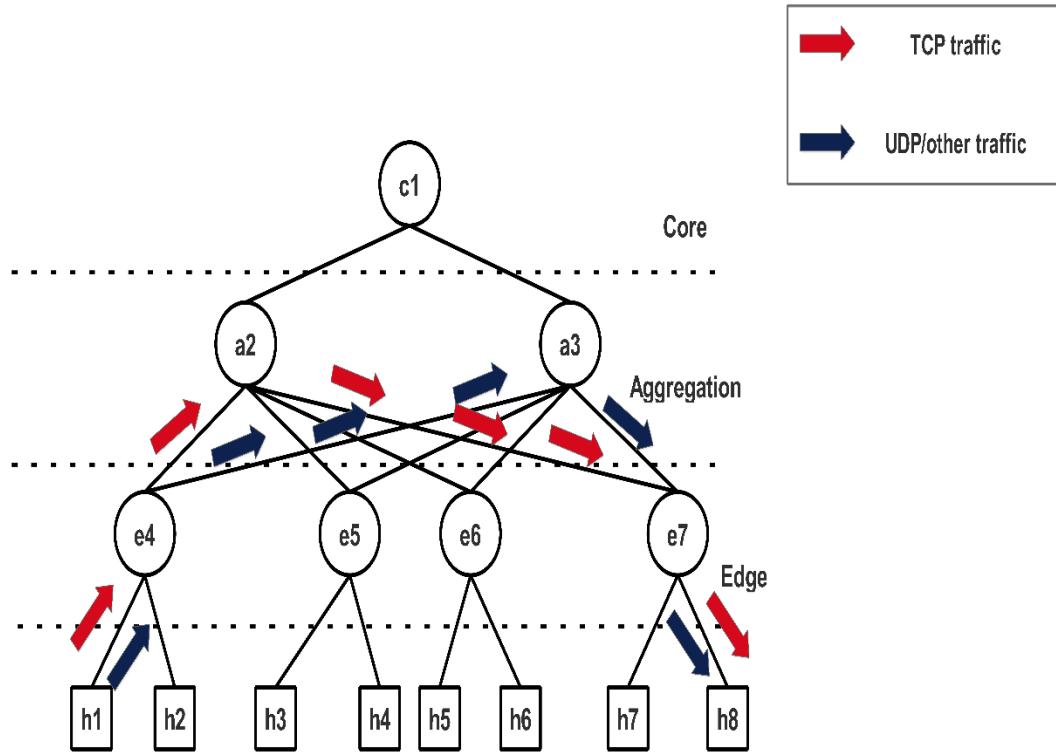
# Shortest Path routing

- Use NetworkX Python package
  - Data center graph representation (G)
  - Evaluate all shortest paths in terms of #hops between all switches (Hint: `all_shortest_paths(G, source, target)`)
  - All this functionality is described in `SwitchWithPaths` class



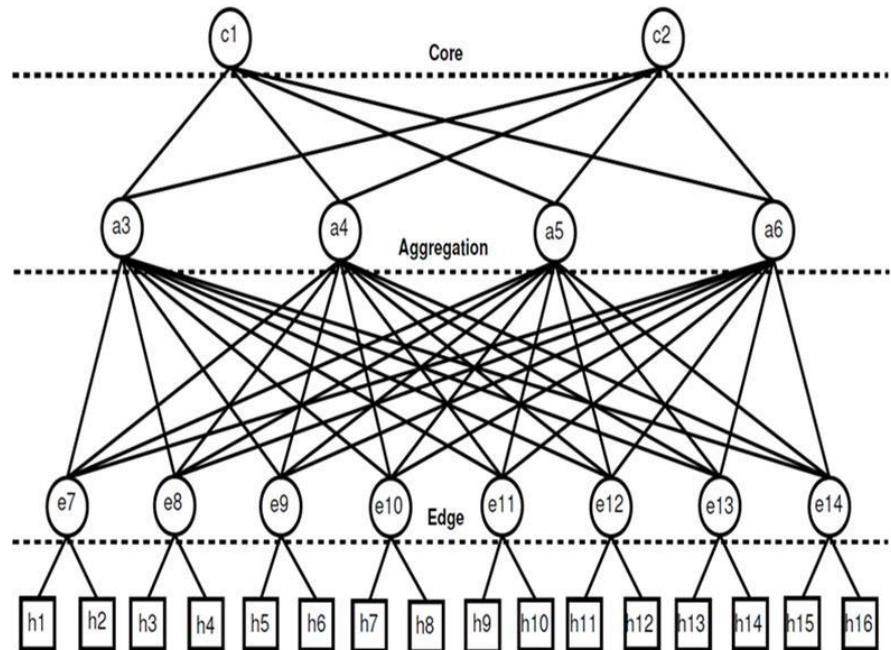
# Traffic Differentiation

- Different type of traffic (TCP vs UDP/other) follow different paths
- Hint: function `getPathsperProto()`



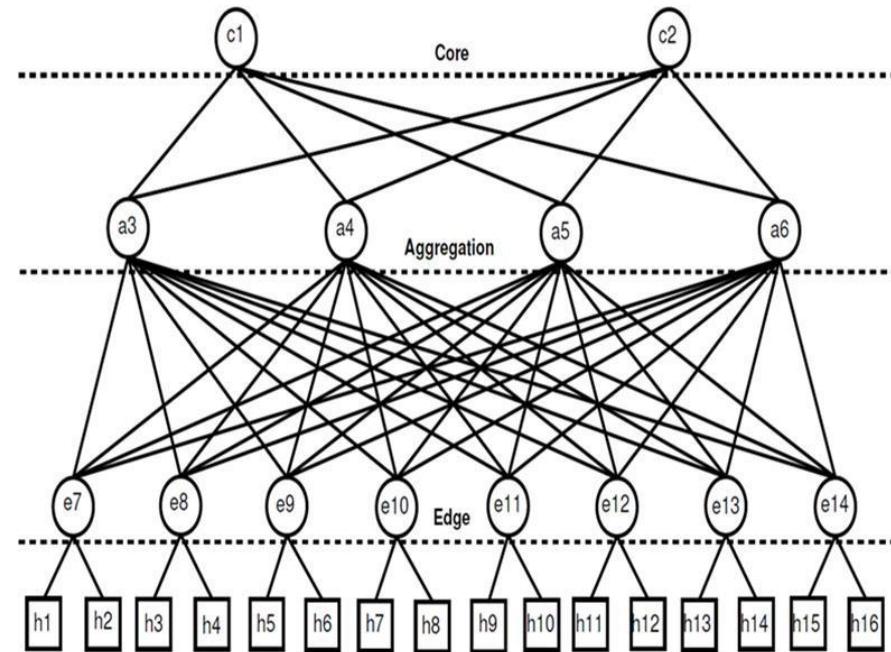
# Flow rules installation

- Flow rule installation is described in `install_end_to_end_IP_path` function
- Called in `PacketIn`'s case
- 2 cases:
  - If `Event switch==dest switch` => install flow rule + packet out
  - Otherwise => path evaluation



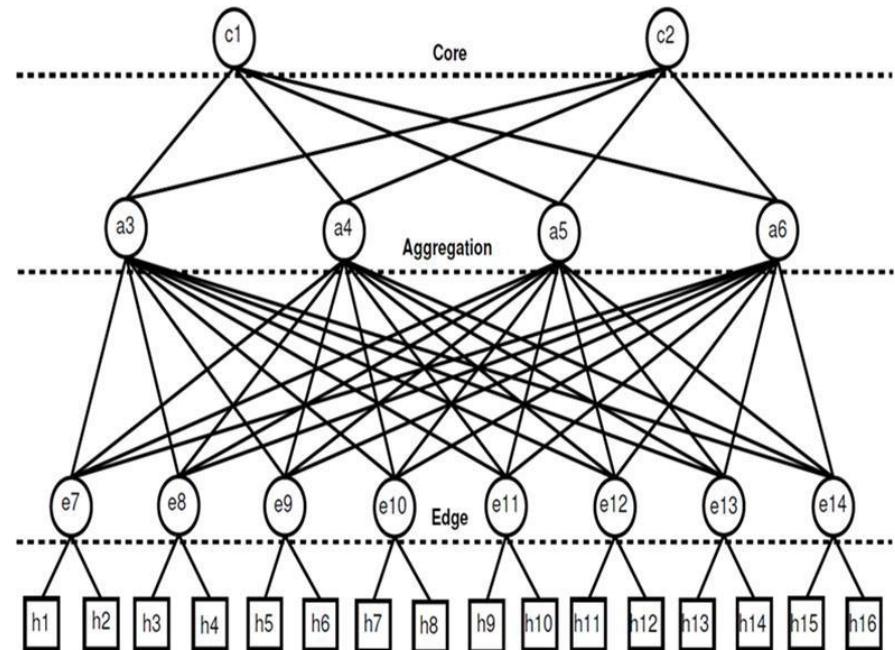
# Flow rules installation

- Topology discovery via POX's Discovery component
- After path evaluation, choose the appropriate path
- Path installation in inverse fashion but in straight direction
- Example: Path: A->B->C->D
  1. C->D
  2. B->C
  3. A->B



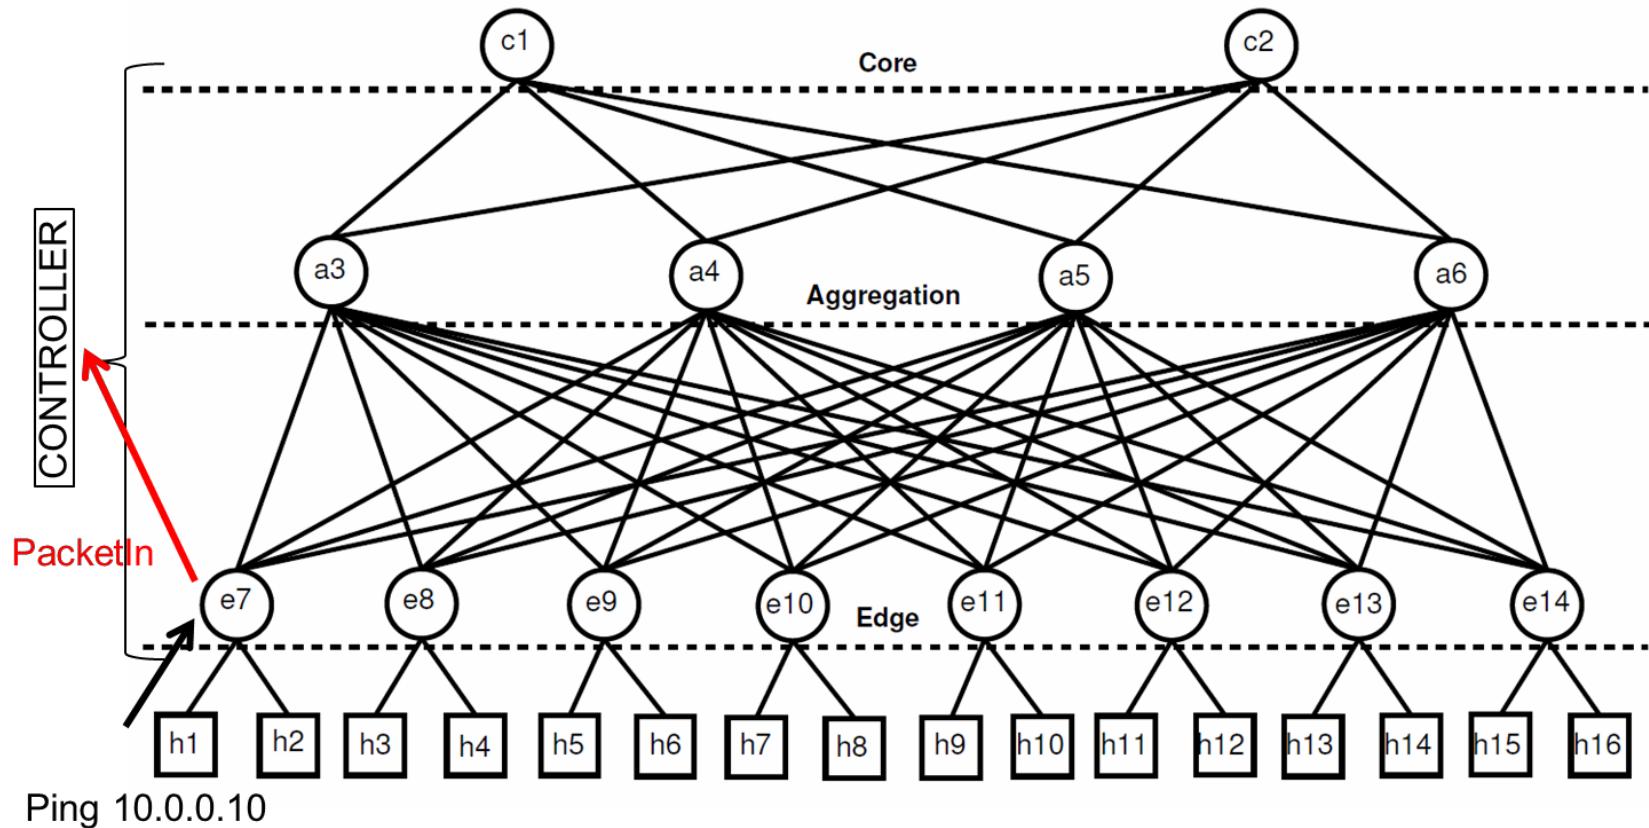
# Flow rules installation

- Flow rule installation via `install_output_flow_rule` function
- Entire packet forwarding (no buffer id used)
- Information about ports interconnection between switches in `sw_sw_ports` variable



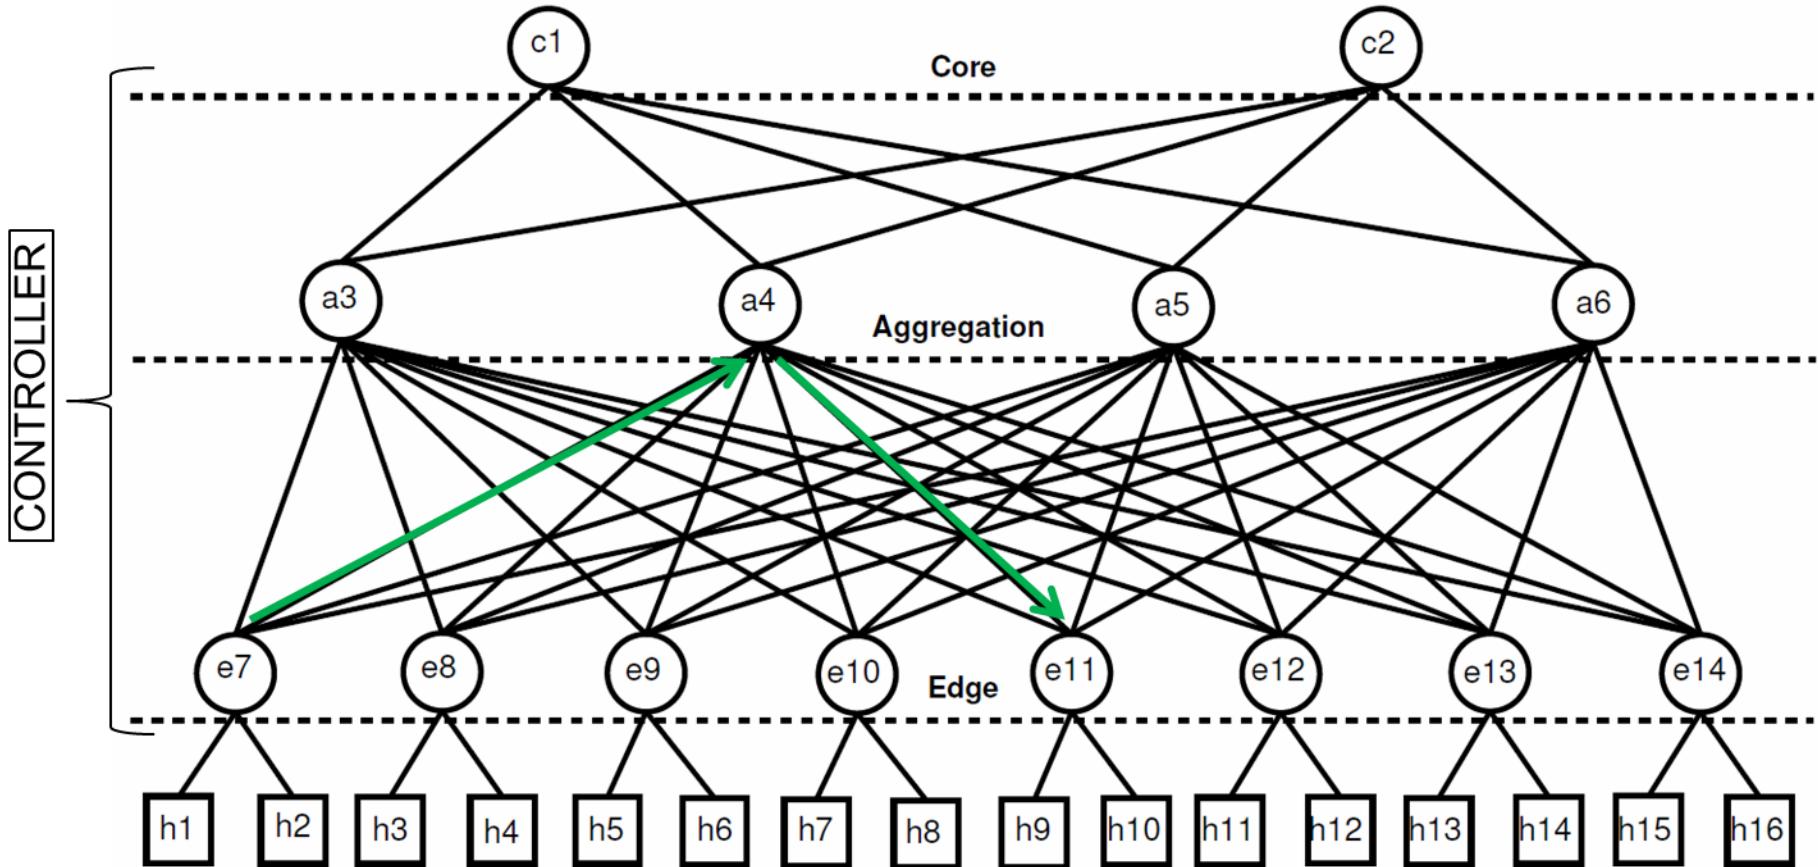
# Flow rules installation - Example

- IP PacketIn



# Flow rules installation - Example

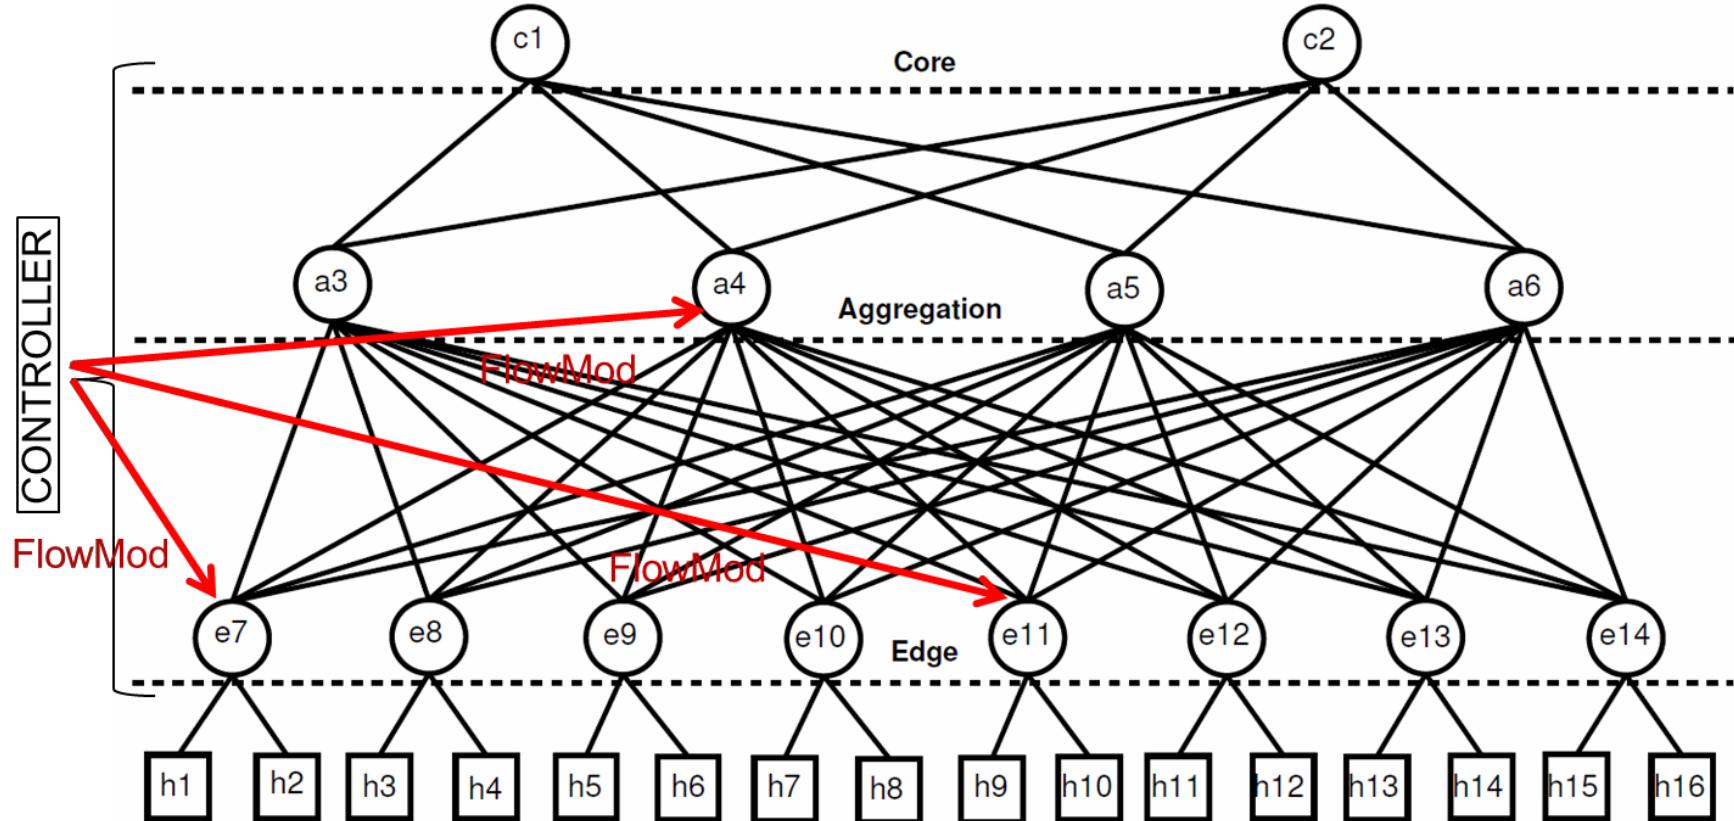
- Path Evaluation



→ Now the controller chooses a random shortest path between  $e_7$  and  $e_{11}$ !

# Flow rules installation - Example

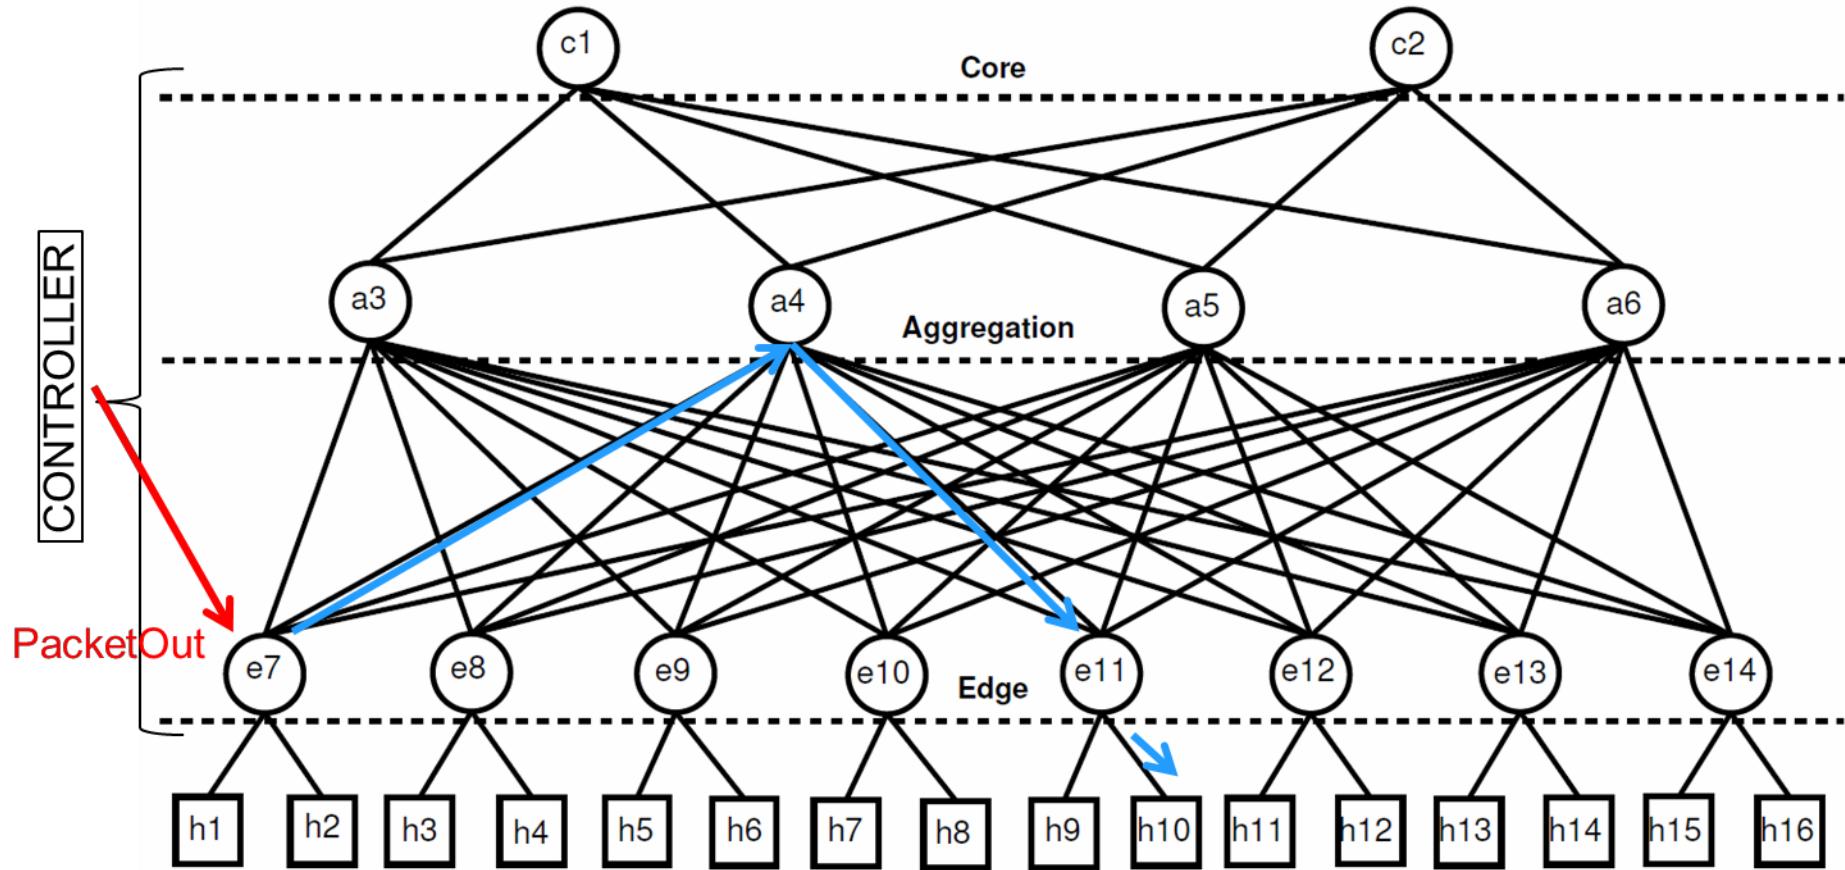
- Path Setup



- Side-note: install flow mods in reverse to avoid even more PacketIns!
- What are potential timing issues here? Can we tackle them?

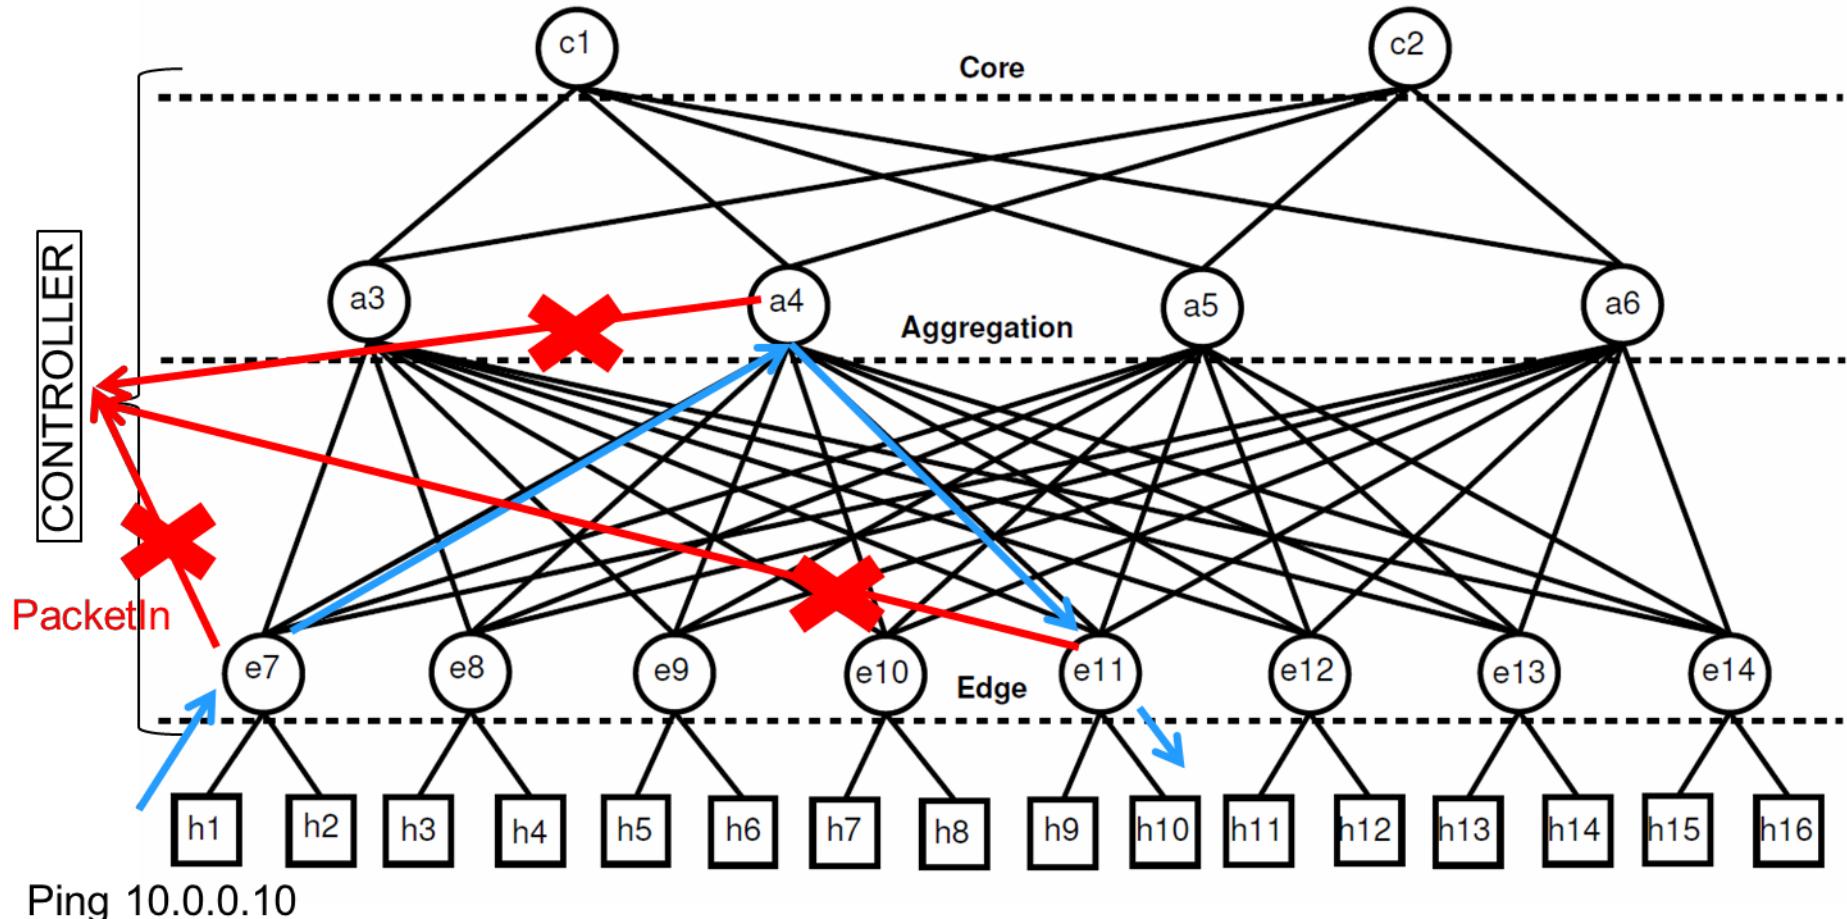
# Flow rules installation - Example

- Current Packet Forwarding



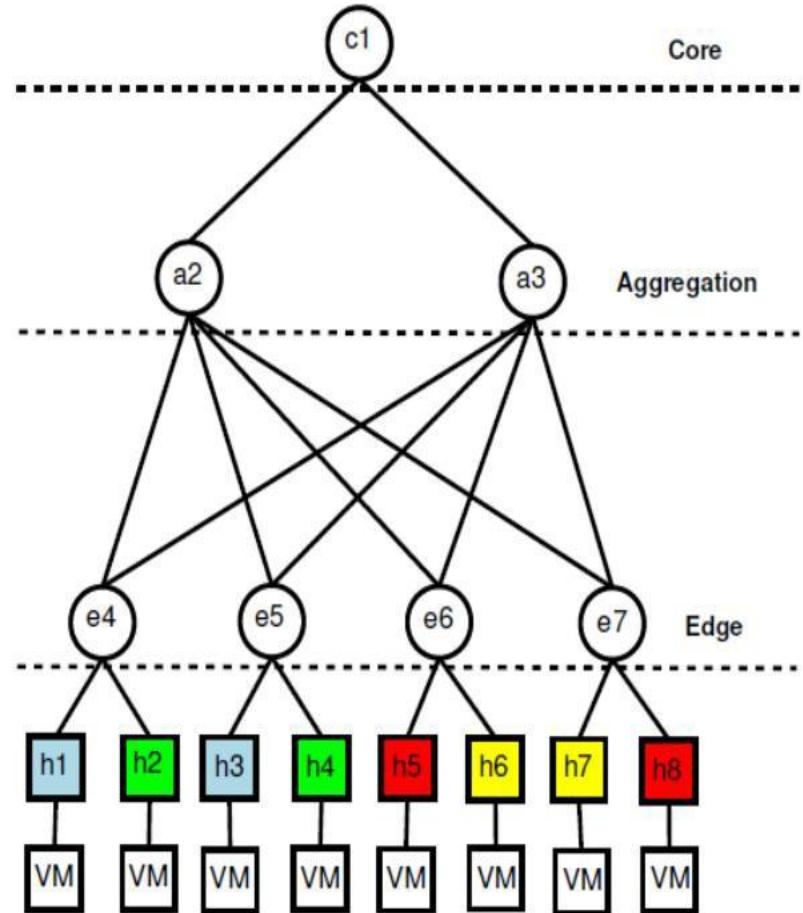
# Flow rules installation - Example

- Coming Packets Forwarding



# Step 3: Firewall/Isolation

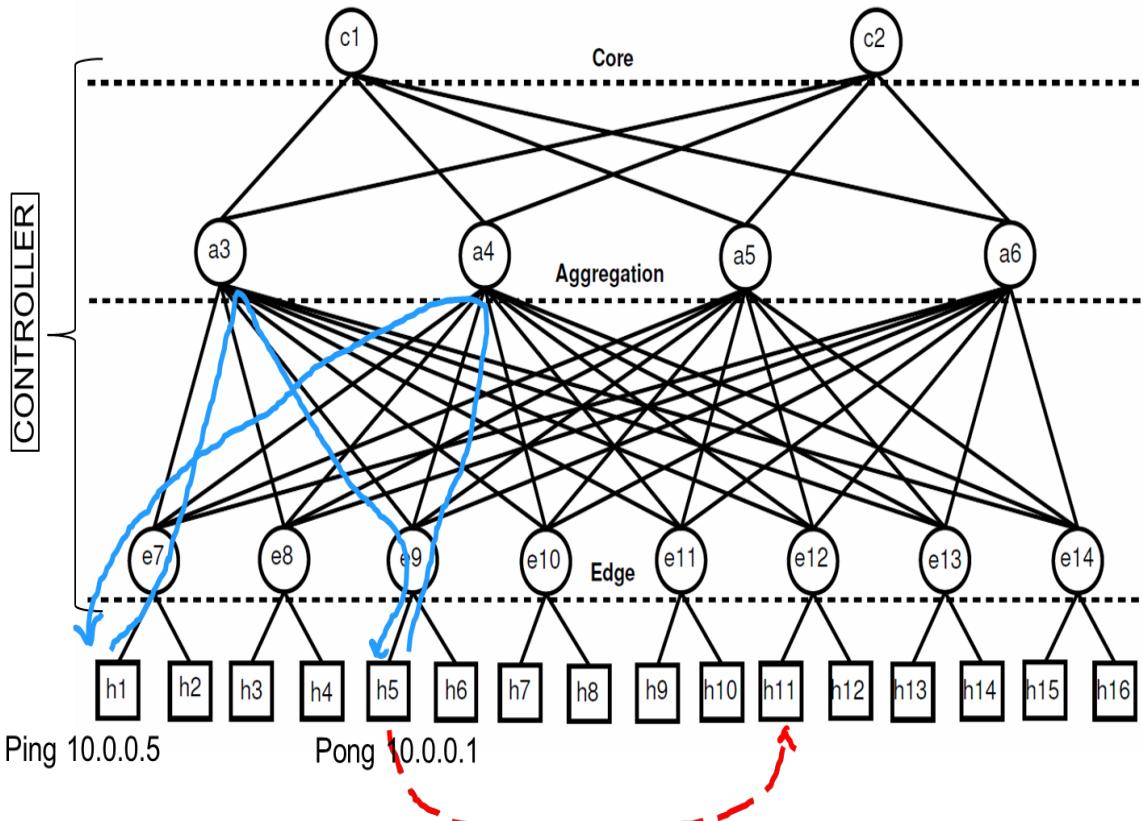
- Communication is allowed between hosts that belong to the same tenant group(same color)
- Firewall policies described in `firewall_policies.csv`
- Policies refer to **both IP and ARP traffic**



# Step 4: Migration/request

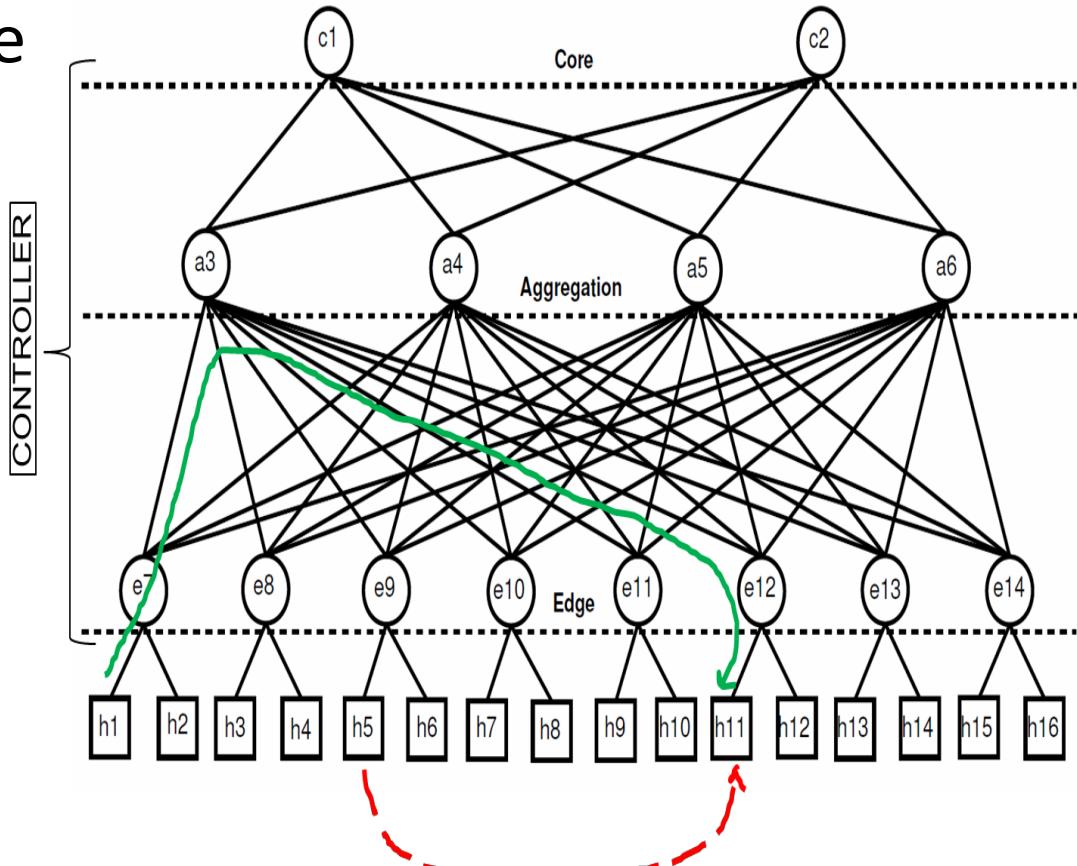
- Migration reason:
  - Maintenance
  - Failure
  - Resource allocation optimization
- Objective:  
**Transparency in IP level**
- Migration event described in  
migration\_events.csv

**Example: h1 ping h5 which is migrated to h11**



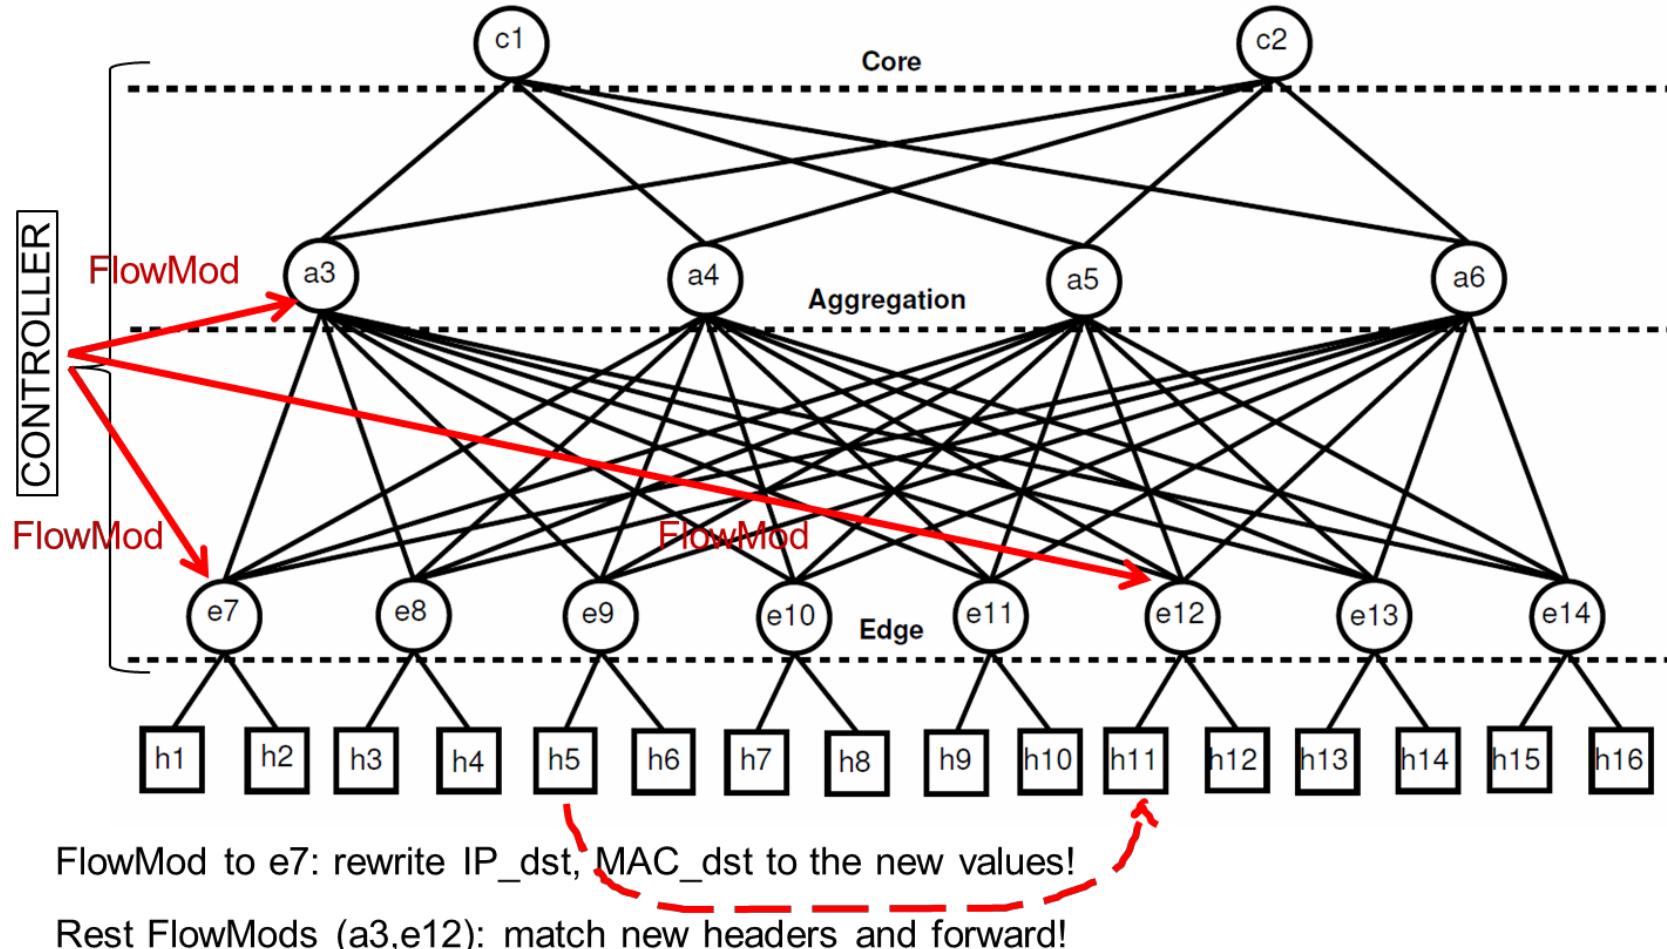
# Step 4: Migration/request

- After migration, all installed flow rules are removed
- New PacketIn events are raised by packets destined to migrated host
- Path evaluation towards h11
- Destination IP and MAC re-write required



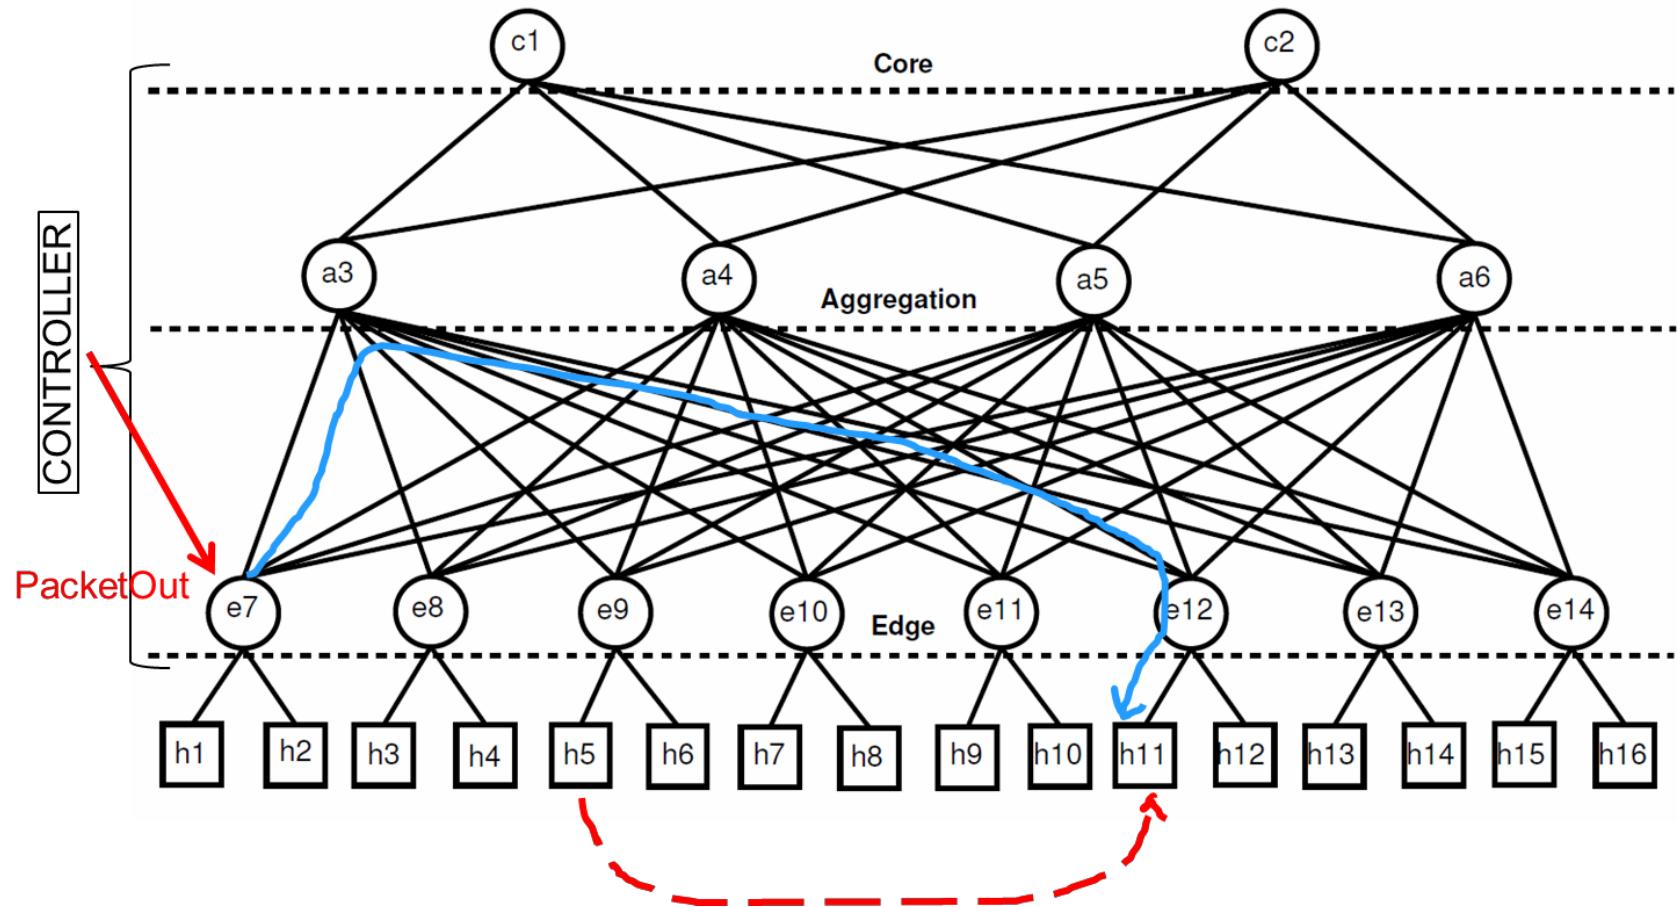
# Step 4: Migration/request

- New flow rules installation



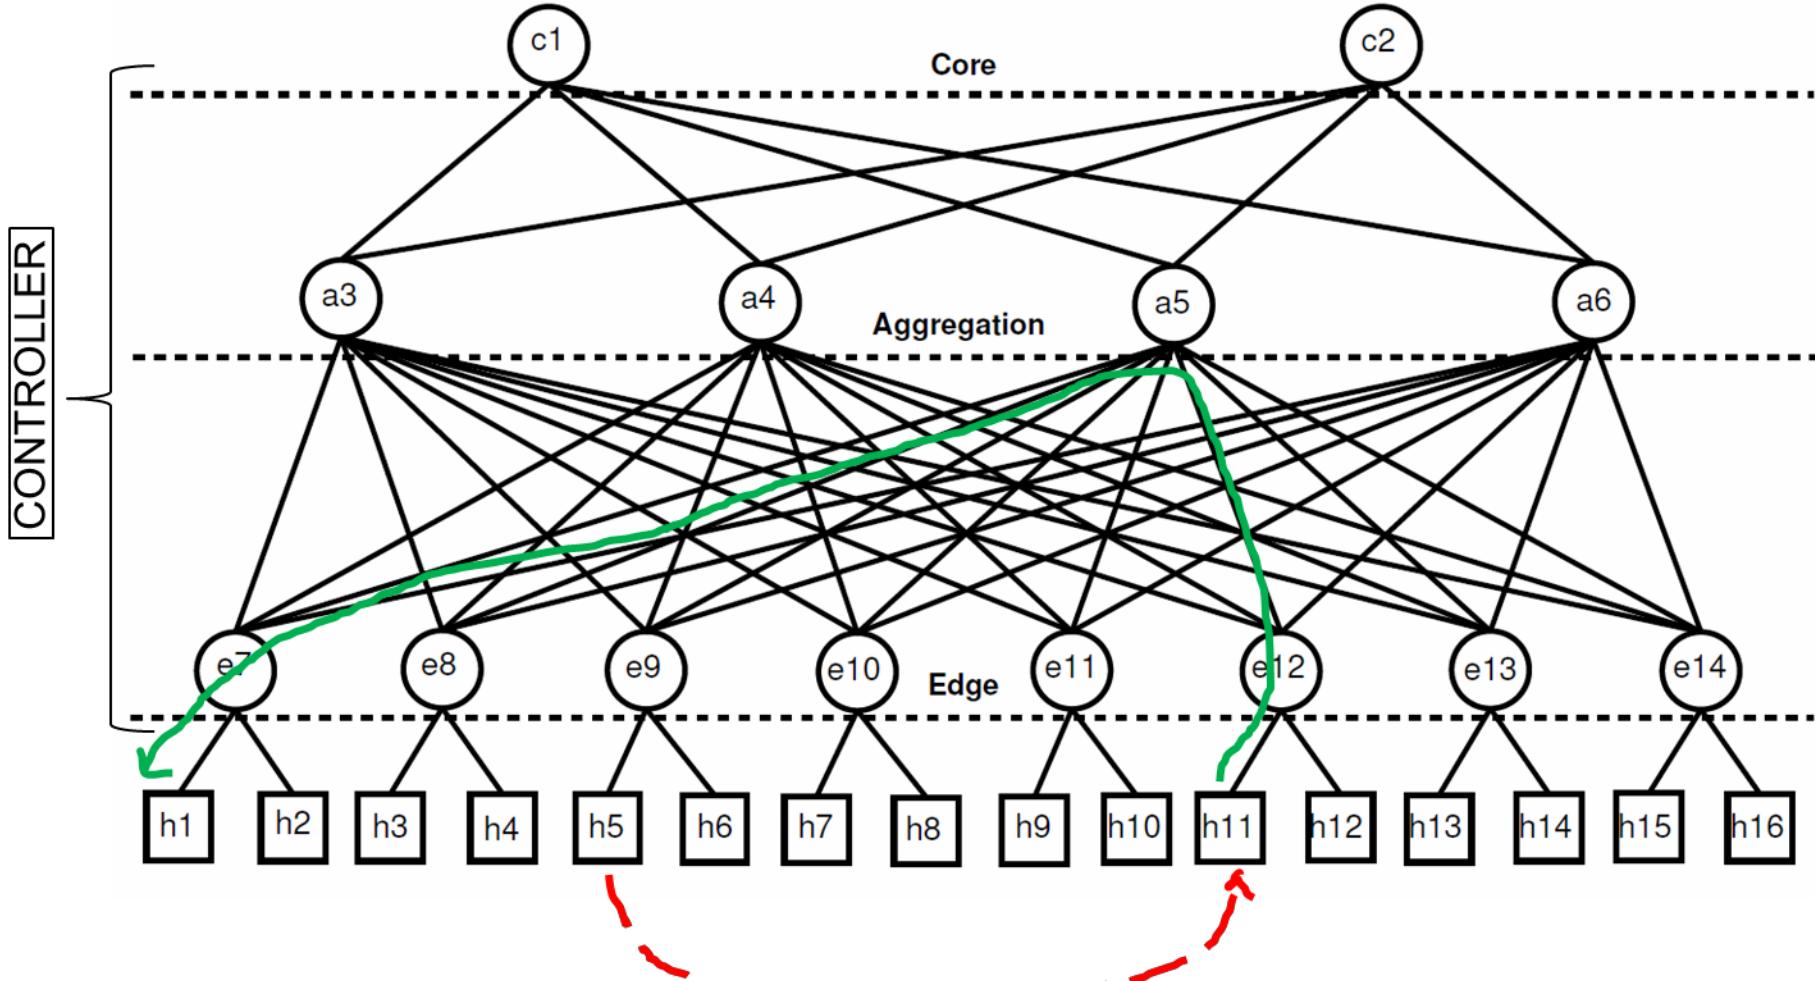
# Step 4: Migration/request

- Packet processing/forwarding



# Step 4: Migration/reply

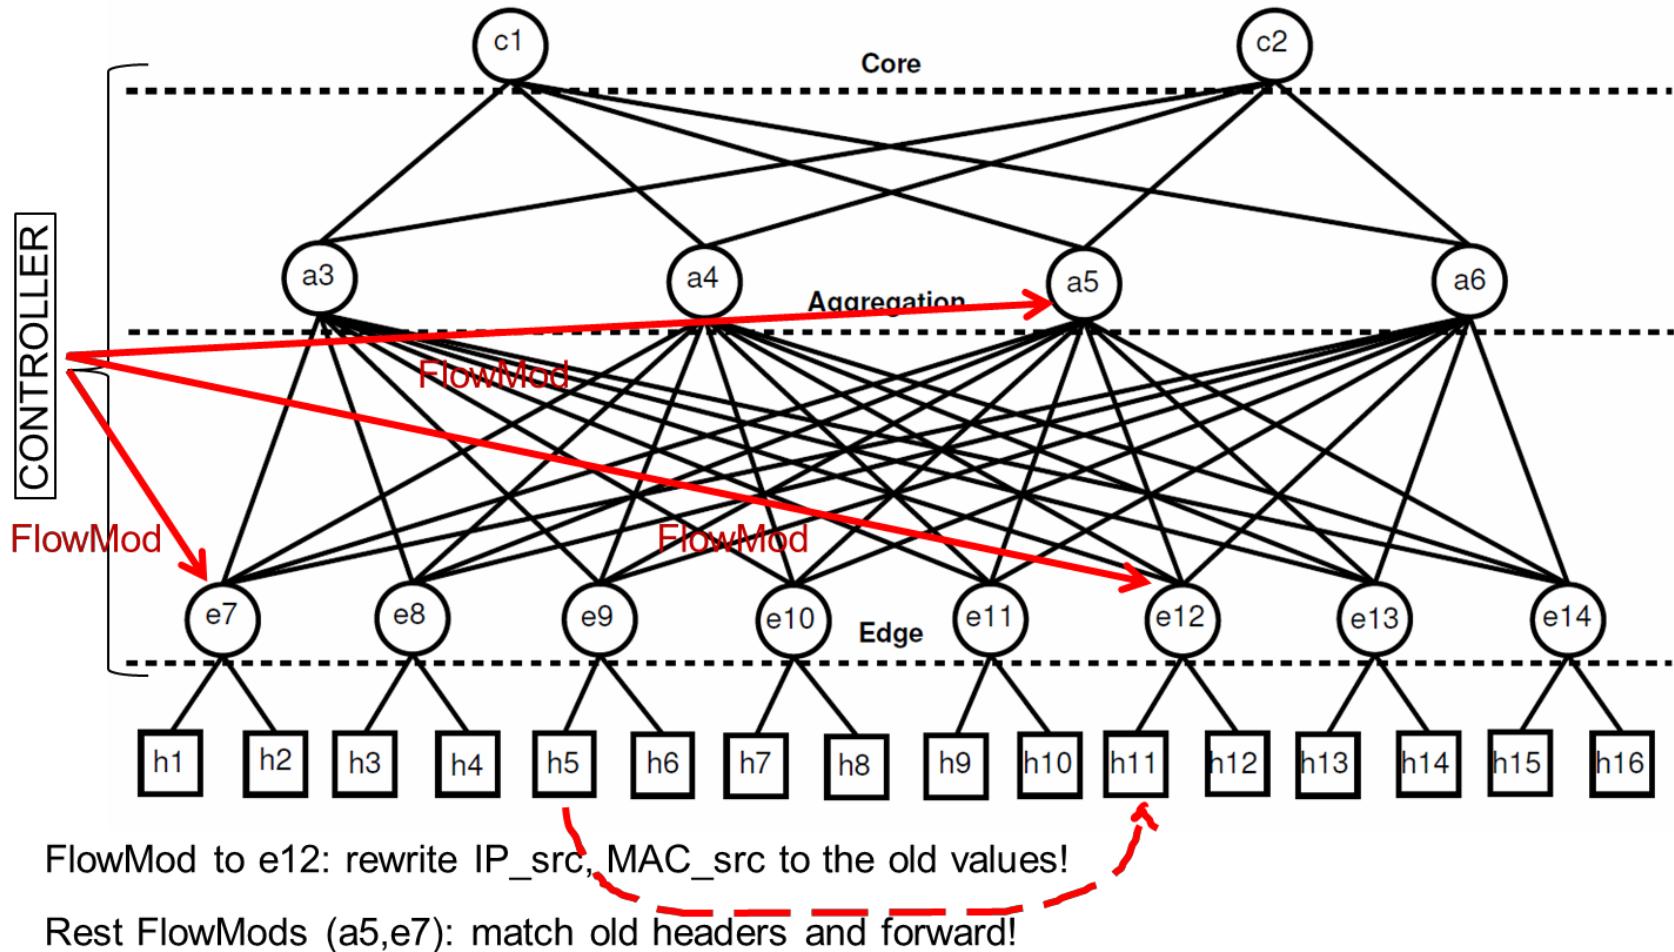
- Reply Path evaluation



# Step 4: Migration/reply

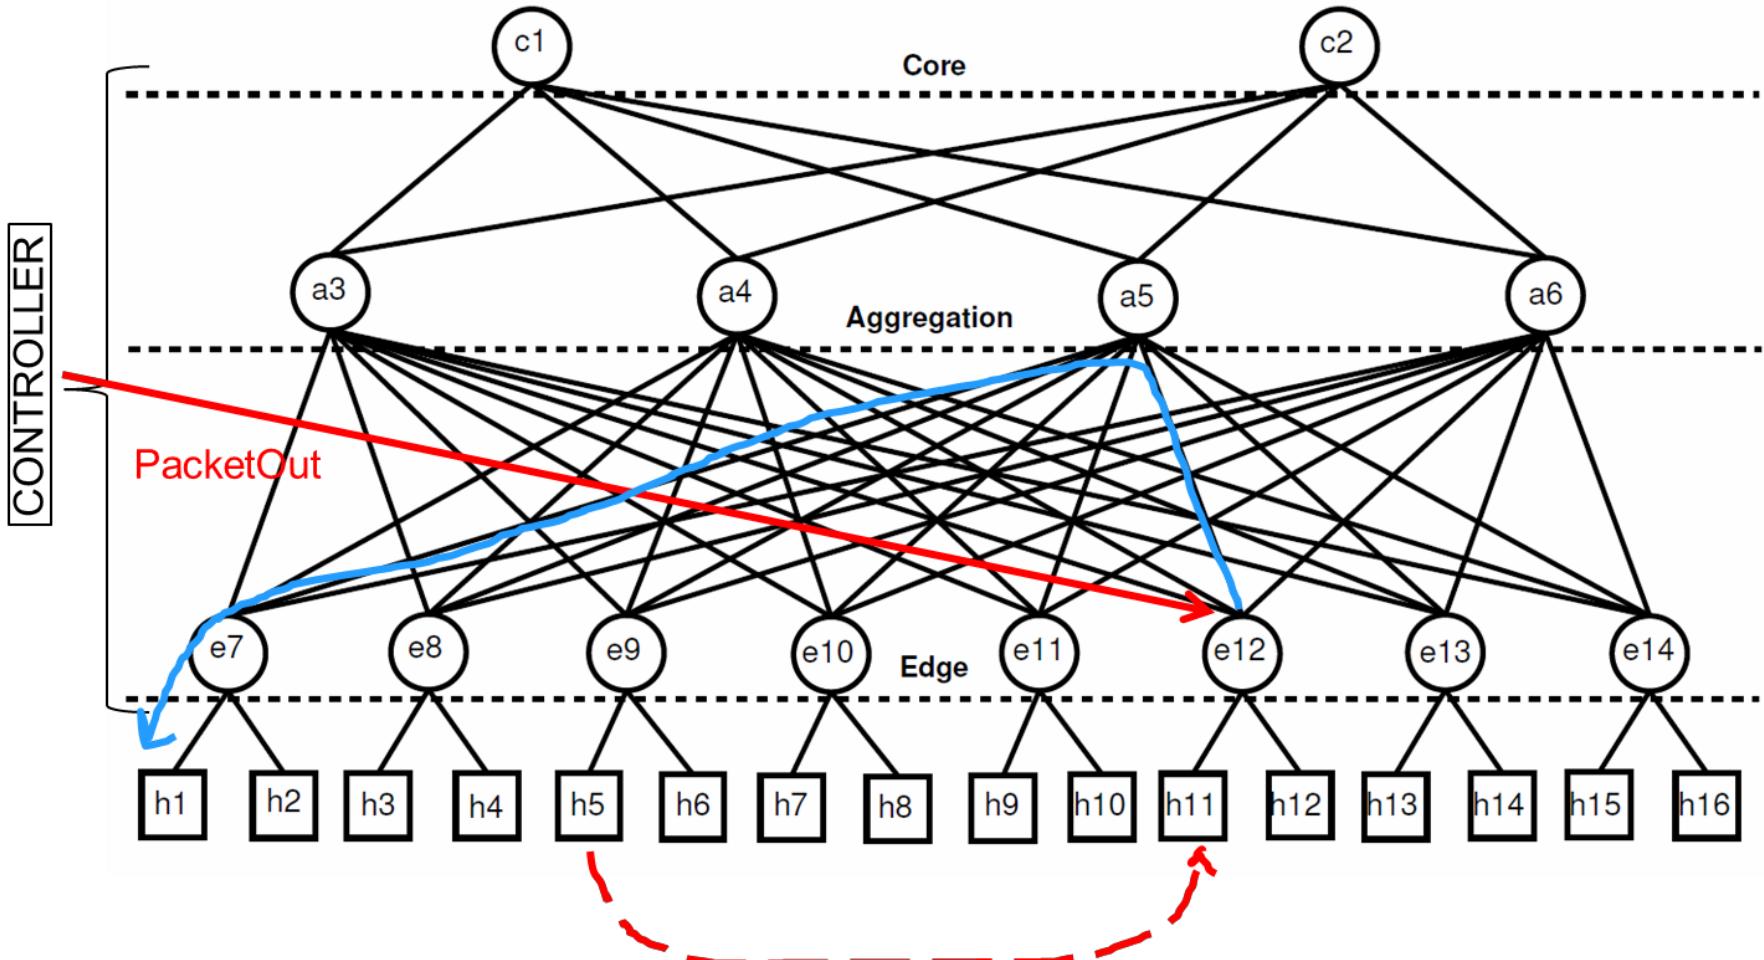
- Reply Path Setup

Source IP and MAC re-write required



# Step 4: Migration/reply

- Reply processing/forwarding



# General guidelines

- Write your code in
  - clos\_topo.py
  - CloudNetController.py
- Read carefully the assignment description
- For your questions, use Moodle forum
- 3 exercise sessions devoted to assignment 2
- **Good luck**

